

Part 3 - Forecasting Regimes

June 24, 2020

1 Module 5: Regime Prediction with Machine Learning - Part 3

In this part we are going to predict economic regimes with machine learning algorithms using Python's Scikit Learn package.

1.1 Table of Contents:

1. Section ??
2. Section ??
3. Section ??
4. Section ??
5. Section ??

1.2 1. Set Up Environment and Read Data

```
In [1]: #load libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn import metrics

from sklearn.metrics import classification_report
from sklearn.metrics import confusion_matrix
from sklearn.metrics import f1_score
from sklearn.metrics import roc_auc_score, roc_curve, auc
from sklearn import model_selection
from sklearn import preprocessing
from sklearn.model_selection import TimeSeriesSplit
from sklearn.feature_selection import SelectFromModel

from sklearn.linear_model import LogisticRegression
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.ensemble import AdaBoostClassifier
from sklearn.ensemble import RandomForestClassifier
```

```
from sklearn.tree import DecisionTreeClassifier
import xgboost as xgb
```

```
from matplotlib import pyplot as mp
import seaborn as sns
import os
import warnings
warnings.filterwarnings('ignore')
```

```
In [2]: df=pd.read_csv('Dataset_Cleaned.csv')
```

```
Label = df["Regime"].apply(lambda regime: 1. if regime == 'Normal' else 0.)
df.insert(loc=2, column="Label", value=Label.values)
```

```
df.head()
```

```
Out[2]:
```

	Date	Regime	Label	RPI	W875RX1	DPCERA3M086SBEA	CMRMTSPLx \
0	9/1/60	Recession	0.0	-0.277761	-0.295013	-0.212118	0.401363
1	10/1/60	Recession	0.0	-0.204460	-0.161434	-0.204519	-0.978339
2	11/1/60	Recession	0.0	-0.520186	-0.577658	-0.785798	-0.611700
3	12/1/60	Recession	0.0	-0.531320	-0.581530	-1.397470	0.071269
4	1/1/61	Normal	1.0	0.054513	0.096044	-0.375484	-1.608684

	RETAILx	INDPRO	IPFPNSS	...	DTCTHFMN 3M lag	DTCTHFMN 6M lag \
0	-0.321981	-0.764446	-0.632790	...	-0.076853	-0.092745
1	-0.239615	-0.308591	0.059474	...	-0.081327	-0.084890
2	-0.435990	-0.935469	-0.748034	...	-0.084509	-0.087674
3	-0.375495	-1.163499	-0.921200	...	-0.095048	-0.076637
4	-0.307555	-0.194576	-0.286457	...	-0.103002	-0.081112

	DTCTHFMN 9M lag	DTCTHFMN 12M lag	DTCTHFMN 18M lag	INVEST 3M lag \
0	-0.087114	-0.085455	-0.099002	-0.454463
1	-0.103429	-0.089335	-0.090643	-0.255545
2	-0.095868	-0.095503	-0.087956	-0.404823
3	-0.094476	-0.087146	-0.076214	-0.331993
4	-0.086616	-0.103462	-0.076910	-0.250758

	INVEST 6M lag	INVEST 9M lag	INVEST 12M lag	INVEST 18M lag
0	-0.566992	-0.395095	-0.450998	-0.089005
1	-0.415804	-0.483302	-0.473520	0.147174
2	-0.392308	-0.605015	-0.445070	-0.091651
3	-0.453979	-0.578110	-0.389378	-0.070555
4	-0.253854	-0.425585	-0.477942	0.167308

```
[5 rows x 711 columns]
```

```
In [3]: df.tail()
```

```
Out[3]:
```

	Date	Regime	Label	RPI	W875RX1	DPCERA3M086SBEA	CMRMTSPLx \
692	5/1/18	Normal	1.0	-0.158774	-0.087869	0.961837	0.428800

693	6/1/18	Normal	1.0	0.349410	0.372881		0.885853	0.063037
694	7/1/18	Normal	1.0	0.225752	0.198647		0.954238	0.150980
695	8/1/18	Normal	1.0	0.496621	0.504524		0.802270	0.788034
696	9/1/18	Normal	1.0	-0.077029	-0.231128		-0.561645	0.393889

	RETAILx	INDPRO	IPFPNSS	...	DTCTHFM 3M lag	DTCTHFM 6M lag	\
692	2.458364	-2.071505	-3.897094	...	-0.482484	0.035214	
693	0.230187	1.119480	1.231754	...	-0.658796	-0.179194	
694	1.059808	0.675562	-0.051761	...	0.108564	-0.311340	
695	-0.474700	1.659097	0.954768	...	0.104177	-0.482299	
696	-0.817945	0.153438	0.379952	...	-0.324269	-0.658623	

	DTCTHFM 9M lag	DTCTHFM 12M lag	DTCTHFM 18M lag	INVEST 3M lag	\
692	-0.213850	0.107884	-0.230363	-2.969053	
693	-0.133821	0.236277	-0.197709	-1.758884	
694	-0.144057	-0.120563	-0.293692	-1.727882	
695	0.033552	-0.213885	-0.613278	0.143563	
696	-0.180970	-0.133855	-0.085913	2.453857	

	INVEST 6M lag	INVEST 9M lag	INVEST 12M lag	INVEST 18M lag
692	0.961673	0.029818	0.229887	-2.082950
693	2.236719	0.274491	-1.001393	-0.345305
694	-0.918590	0.481157	-0.256506	1.560538
695	-2.983834	0.964076	0.037260	0.000571
696	-1.766319	2.250399	0.282925	-0.081374

[5 rows x 711 columns]

1.3 2. Methodology

Our exercise will be based on classification problem. We have two binary outcomes that we want to predict with certain variables. Here we will summarize our approach to predict recessions with machine learning algorithms.

1. We will perform feature selection before making our forecasts. We will use L_1 regularized logistic regression for that purpose.
2. Separate dataset into training and validation datasets. Split based dataset based on time: the period over 1960-1996 is selected for training and the period over 1996-2018 is kept for validation
3. Evaluate performances of the machine learning algorithms on training dataset with cross validation (CV). Since we have time series structure we will use a special type of CV function in Python, `TimeSeriesSplit`. We will use Receiver operating characteristic (ROC) as scoring metric in our models. Related Python functions for this metric are `roc_auc_score` and `roc_curve`.
4. Select the best performing models based on average accuracy and standard deviation of the CV results. We will take logistic regression as a benchmark model since this is the traditional method has been used to approach this problem.

5. Then we make predictions on the validation dataset with selected models. First, we use `GridSearchCV` for selected model on training dataset to find best combination of parameters for the model. Then we evaluate the model on validation dataset and report accuracy metrics and feature importance results.

1.4 3. Feature Selection with L_1 Penalty

```
In [4]: # Time Series Split
df_idx = df[df.Date == '12/1/96'].index[0]

df_targets=df['Label'].values
df_features=df.drop(['Regime','Date','Label'], axis=1)

df_training_features = df.iloc[:df_idx,:].drop(['Regime','Date','Label'], axis=1)
df_validation_features = df.iloc[df_idx:, :].drop(['Regime','Date','Label'], axis=1)

df_training_targets = df['Label'].values
df_training_targets=df_training_targets[:df_idx]

df_validation_targets = df['Label'].values
df_validation_targets=df_validation_targets[df_idx:]

In [5]: print(len(df_training_features),len(df_training_targets),len(df_targets))
        print(len(df_validation_features),len(df_validation_targets),len(df_features))

435 435 697
262 262 697

In [6]: scoring="roc_auc"
kfold= model_selection.TimeSeriesSplit(n_splits=3)
seed=8

# Create regularization hyperparameter space
C = np.reciprocal([0.00000001, 0.00000005, 0.0000001, 0.0000005, 0.000001, 0.000005, 0.00001, 0.0001, 0.0005, 0.001, 0.005, 0.01, 0.05, 0.1, 0.5, 1, 5, 10, 100, 1000])

# Create hyperparameter options
hyperparameters = dict(C=C)

model=LogisticRegression(max_iter=10000,penalty='l1')
LR_penalty=model_selection.GridSearchCV(estimator=model, param_grid= hyperparameters,
                                       cv=kfold, scoring=scoring).fit(X=df_features,
                                                                    y=df_targets).best_estimator_

LR_penalty

Out[6]: LogisticRegression(C=0.1, class_weight=None, dual=False, fit_intercept=True,
                          intercept_scaling=1, max_iter=10000, multi_class='warn',
```

```
n_jobs=None, penalty='l1', random_state=None, solver='warn',
tol=0.0001, verbose=0, warm_start=False)
```

```
In [7]: X=df_features
        y=df_targets
        lr_l1 = LogisticRegression(C=0.1, max_iter=10000,penalty="l1").fit(X,y)
        model = SelectFromModel(lr_l1,prefit=True)
        feature_idx = model.get_support()
        feature_name = X.columns[feature_idx]
        X_new = model.transform(X)
        X_new.shape
```

```
Out[7]: (697, 35)
```

```
In [8]: feature_name
```

```
Out[8]: Index(['DPCERA3M086SBEA', 'IPFPNSS', 'IPMANSICS', 'HWI', 'UNRATE', 'PAYEMS',
               'SRVPRD', 'USTPU', 'USWTRADE', 'USFIRE', 'HOUST', 'HOUSTW',
               'S&P div yield', 'TB3MS', 'TB3SMFFM', 'IPMAT 9M lag', 'IPDMAT 12M lag',
               'CLAIMSx 9M lag', 'DMANEMP 12M lag', 'USFIRE 3M lag', 'AWOTMAN 18M lag',
               'BUSLOANS 6M lag', 'S&P 500 6M lag', 'S&P div yield 3M lag',
               'TB3MS 6M lag', 'AAA 3M lag', 'TB3SMFFM 3M lag', 'T1YFFM 6M lag',
               'T5YFFM 6M lag', 'T10YFFM 6M lag', 'T10YFFM 12M lag', 'EXUSUKx 12M lag',
               'WPSID62 9M lag', 'OILPRICEx 9M lag', 'CUSR0000SAD 3M lag'],
              dtype='object')
```

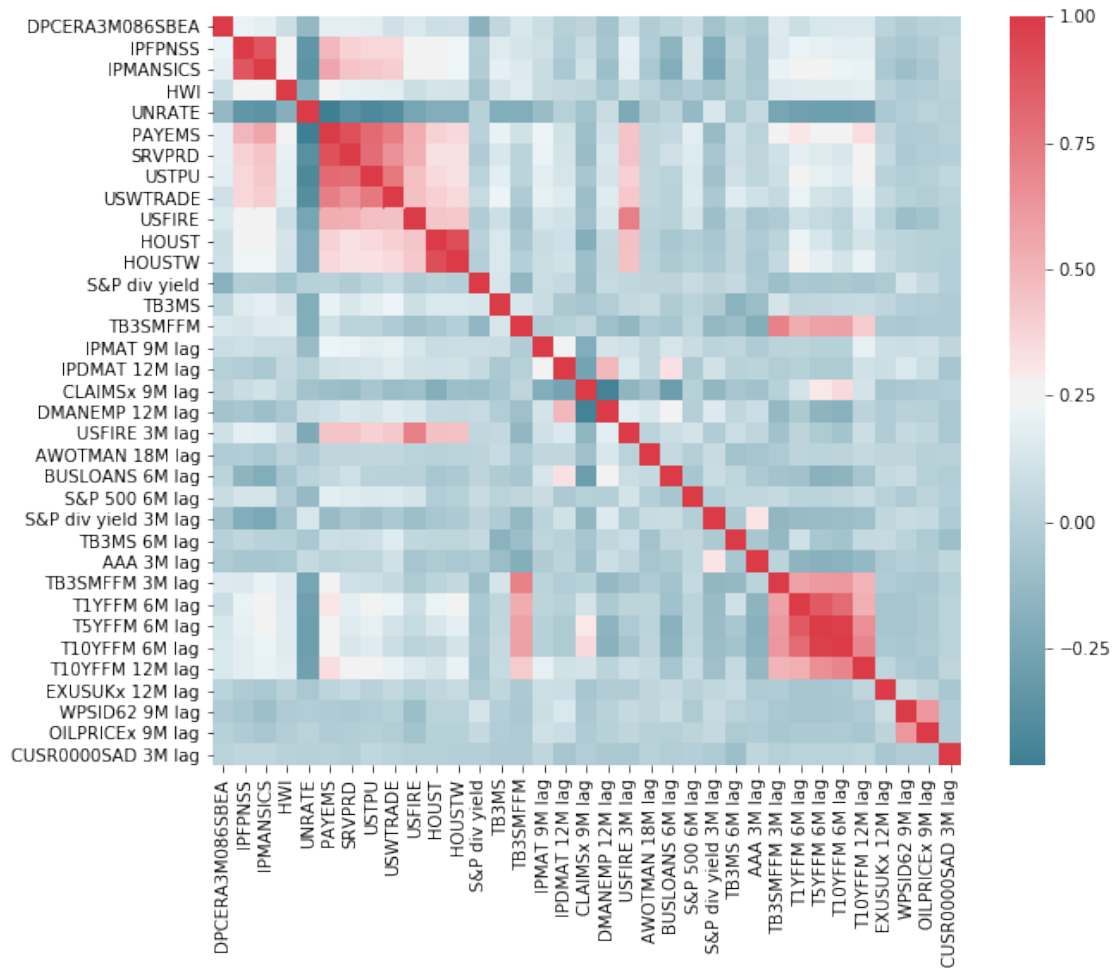
```
In [9]: df_2=df[feature_name]
        df_2.insert(loc=0, column="Date", value=df['Date'].values)
        df_2.insert(loc=1, column="Regime", value=df['Regime'].values)
        df_2.insert(loc=2, column="Label", value=df['Label'].values)

        df_2.head()
        df_2.shape
```

```
Out[9]: (697, 38)
```

```
In [10]: corr = df_2.drop(['Date', 'Regime', 'Label'],axis=1).corr()
         plt.figure(figsize=(10, 8))
         sns.heatmap(corr, mask=np.zeros_like(corr, dtype=np.bool),
                     cmap=sns.diverging_palette(220, 10, as_cmap=True), square=True)
```

```
Out[10]: <matplotlib.axes._subplots.AxesSubplot at 0x10e77fb38>
```



1.5 4. Training Algorithms on Training Dataset

For a detail description of the machine learning algorithms you can check scikit-learn's documentation [here](#).

```
In [11]: df=df_2
         df.shape
```

```
Out[11]: (697, 38)
```

```
In [12]: # Time Series Split
         df_idx = df[df.Date == '12/1/96'].index[0]

         df_targets=df['Label'].values
         df_features=df.drop(['Regime', 'Date', 'Label'], axis=1)

         df_training_features = df.iloc[:df_idx,:].drop(['Regime', 'Date', 'Label'], axis=1)
```

```

df_validation_features = df.iloc[df_idx:, :].drop(['Regime', 'Date', 'Label'], axis=1)

df_training_targets = df['Label'].values
df_training_targets=df_training_targets[:df_idx]

df_validation_targets = df['Label'].values
df_validation_targets=df_validation_targets[df_idx:]

```

In [13]: seed=8

```

scoring='roc_auc'
kfold = model_selection.TimeSeriesSplit(n_splits=3)
models = []

models.append(('LR', LogisticRegression(C=1e09)))
models.append(('LR_L1', LogisticRegression(penalty = 'l1')))
models.append(('LR_L2', LogisticRegression(penalty = 'l2')))
models.append(('LDA', LinearDiscriminantAnalysis()))
models.append(('KNN', KNeighborsClassifier()))
models.append(('GB', GradientBoostingClassifier()))
models.append(('ABC', AdaBoostClassifier()))
models.append(('RF', RandomForestClassifier()))
models.append(('XGB', xgb.XGBClassifier()))

results = []
names = []
lb = preprocessing.LabelBinarizer()

for name, model in models:
    cv_results = model_selection.cross_val_score(estimator = model, X = df_training_features,
                                                y = lb.fit_transform(df_training_targets),
                                                cv=kfold)

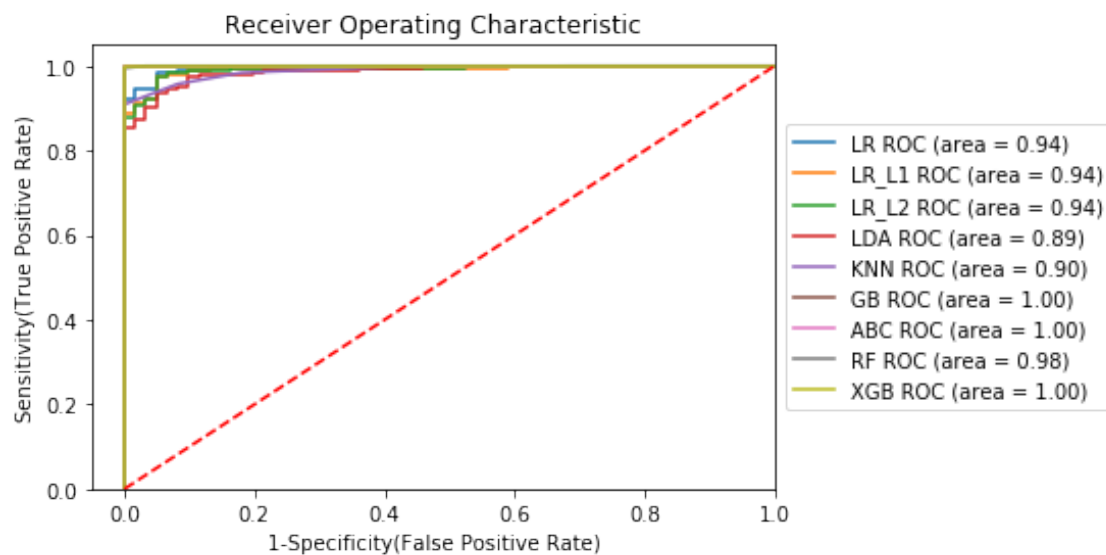
    model.fit(df_training_features, df_training_targets) # train the model
    fpr, tpr, thresholds = metrics.roc_curve(df_training_targets, model.predict_proba(df_training_features),
                                              probabilities=True)
    auc = metrics.roc_auc_score(df_training_targets,model.predict(df_training_features))
    plt.plot(fpr, tpr, label='%s ROC (area = %0.2f)' % (name, auc))
    results.append(cv_results)
    names.append(name)
    msg = "%s: %f (%f)" % (name, cv_results.mean(), cv_results.std())
    print(msg)

plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([-0.05, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('1-Specificity(False Positive Rate)')
plt.ylabel('Sensitivity(True Positive Rate)')
plt.title('Receiver Operating Characteristic')
plt.legend(loc='center left', bbox_to_anchor=(1, 0.5))
plt.show()

```

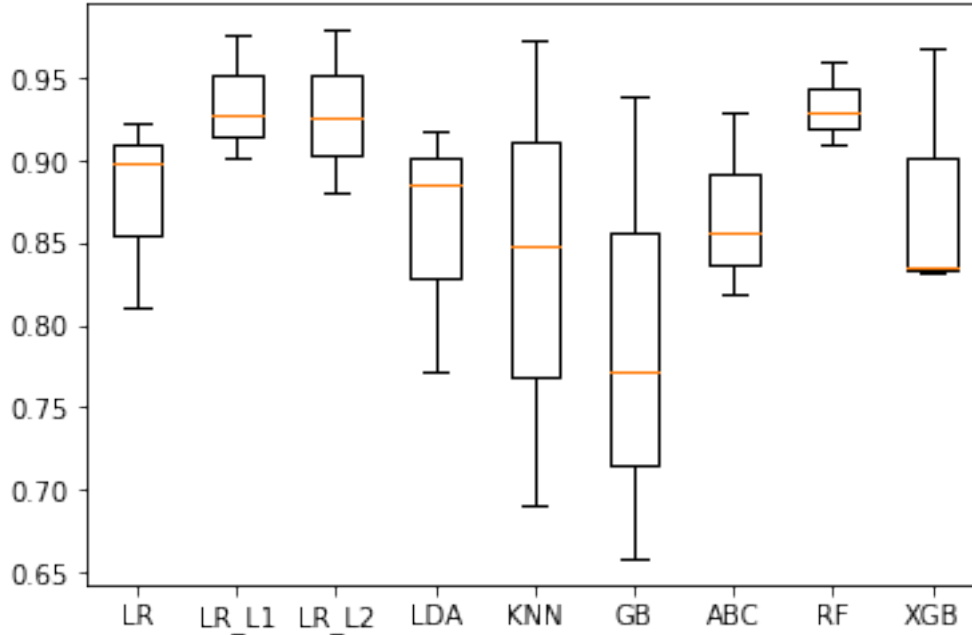
```
warnings.filterwarnings('ignore')
```

```
LR: 0.877570 (0.047838)
LR_L1: 0.934921 (0.031076)
LR_L2: 0.928582 (0.040605)
LDA: 0.858740 (0.063361)
KNN: 0.837252 (0.115853)
GB: 0.790105 (0.115667)
ABC: 0.868133 (0.045829)
RF: 0.933246 (0.020609)
XGB: 0.878555 (0.063837)
```



```
In [14]: fig = plt.figure()
fig.suptitle('Algorithm Comparison based on Cross Validation Scores')
ax = fig.add_subplot(111)
plt.boxplot(results)
ax.set_xticklabels(names)
plt.show()
```


Algorithm Comparison based on Cross Validation Scores



1.6 4. Evaluate Performances of the Algorithms on Validation Dataset

1.6.1 Logistic Regression

Logistic regression is the most commonly used statistical model for binary classification. It uses the logit model of relating log-odds of the dependent variable linearly with the predictor (explanatory) variables to learn a form of the following logistic function that is used to separate instances of the two different classes.

$$Pr(y = 1|x) = h_{\beta}(x) = \frac{1}{1 + e^{-(\beta_0 + \beta_1 x)}} \text{ where } \beta_0 \text{ is the intercept and } \beta \text{ is the vector of trained weights}$$

The function returns a probability measure of which class a new instance is given its features, this measure is then used to make the final classification with some probability threshold, traditionally being 0.5.

The specific logistic function is learned through gradient descent which focuses on minimizing error calculated through some cost function. A typical approach is to use the following maximum-likelihood estimator to minimize error in predicted probabilities.

$$J(\beta) = -\frac{1}{N} \sum_{i=1}^N y_i \log(h_{\beta}(x_i)) + (1 - y_i) \log(1 - h_{\beta}(x_i))$$

where N is the number of observations, $h_{\beta}(x)$ is as defined above, and y_i is the predicted class

A regularization term

$$\lambda \sum_{i=1}^N |\beta_i| \text{ where } \lambda \text{ is a regularization parameter,}$$

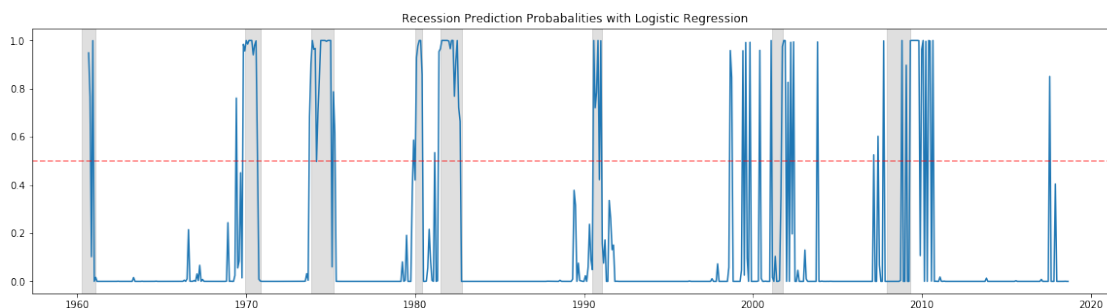
is often added to the cost function to prevent overfitting by penalizing large coefficients. This regularization can be L1 or L2 penalty depending on the problem at hand.

```
In [15]: model=LogisticRegression(C=1e09) # high penalty
LR=model.fit(df_training_features,df_training_targets)
training_predictions=LR.predict(df_training_features)
prob_predictions = LR.predict_proba(df_training_features)
prob_predictions = np.append(prob_predictions, LR.predict_proba(df_validation_features))
```

```
In [16]: import datetime
# define periods of recession
rec_spans = []
#rec_spans.append([datetime.datetime(1957,8,1), datetime.datetime(1958,4,1)])
rec_spans.append([datetime.datetime(1960,4,1), datetime.datetime(1961,2,1)])
rec_spans.append([datetime.datetime(1969,12,1), datetime.datetime(1970,11,1)])
rec_spans.append([datetime.datetime(1973,11,1), datetime.datetime(1975,3,1)])
rec_spans.append([datetime.datetime(1980,1,1), datetime.datetime(1980,6,1)])
rec_spans.append([datetime.datetime(1981,7,1), datetime.datetime(1982,10,1)])
rec_spans.append([datetime.datetime(1990,7,1), datetime.datetime(1991,2,1)])
rec_spans.append([datetime.datetime(2001,3,1), datetime.datetime(2001,10,1)])
rec_spans.append([datetime.datetime(2007,12,1), datetime.datetime(2009,5,1)])
```

```
In [17]: sample_range = pd.date_range(start='9/1/1960', end='9/1/2018', freq='MS')
```

```
plt.figure(figsize=(20,5))
plt.plot(sample_range.to_series().values, prob_predictions[:,0])
for i in range(len(rec_spans)):
    plt.axvspan(rec_spans[i][0], rec_spans[i][len(rec_spans[i]) - 1], alpha=0.25, color='r')
plt.axhline(y=0.5, color='r', ls='dashed', alpha = 0.5)
plt.title('Recession Prediction Probabilities with Logistic Regression')
mp.savefig('plot1.png', bbox_inches='tight')
plt.show()
```



1.6.2 Logistic Regression with Regularization

```
In [18]: # Create regularization penalty space
penalty = ['l1', 'l2']

# Create regularization hyperparameter space
C = np.reciprocal([0.00000001, 0.00000005, 0.0000001, 0.0000005, 0.000001, 0.000005, 0.00001, 0.0001, 0.0005, 0.001, 0.005, 0.01, 0.05, 0.1, 0.5, 1, 5, 10])

# Create hyperparameter options
hyperparameters = dict(C=C, penalty=penalty)

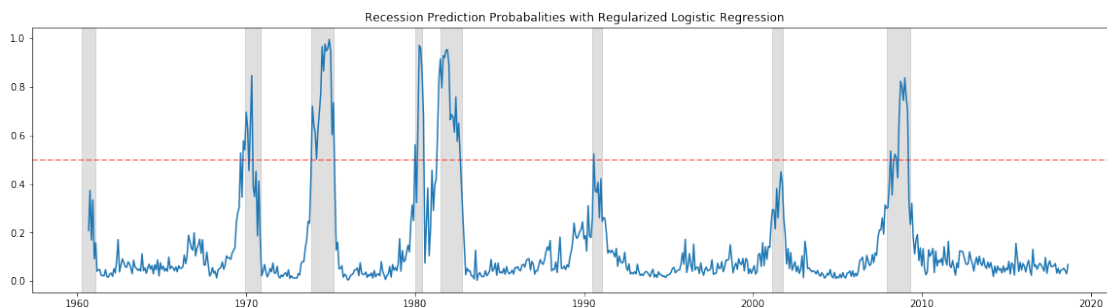
model=LogisticRegression(max_iter=10000)
LR_penalty=model_selection.GridSearchCV(estimator=model, param_grid= hyperparameters,
                                         cv=kfold, scoring=scoring).fit(df_training_features, df_training_targets)

training_predictions=LR_penalty.predict(df_training_features)

In [19]: prob_predictions = LR_penalty.predict_proba(df_training_features)
prob_predictions = np.append(prob_predictions, LR_penalty.predict_proba(df_validation_features), axis=1)

In [20]: sample_range = pd.date_range(start='9/1/1960', end='9/1/2018', freq='MS')

plt.figure(figsize=(20,5))
plt.plot(sample_range.to_series().values, prob_predictions[:,0])
for i in range(len(rec_spans)):
    plt.axvspan(rec_spans[i][0], rec_spans[i][len(rec_spans[i]) - 1], alpha=0.25, color='gray')
plt.axhline(y=0.5, color='r', ls='dashed', alpha = 0.5)
plt.title('Recession Prediction Probabilities with Regularized Logistic Regression')
mp.savefig('plot2.png', bbox_inches='tight')
plt.show()
```



1.6.3 XGBoosting

```
In [21]: xgboost = model_selection.GridSearchCV(estimator=xgb.XGBClassifier(),
                                                param_grid={'booster': ['gbtree']},
```

```

scoring=scoring, cv=kfold).fit(df_training_features, df_training_targets)
lb.fit_transform(df_validation_features)

xgboost.fit(df_training_features, df_training_targets)

Out [21]: XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1,
                        colsample_bytrees=1, gamma=0, learning_rate=0.1, max_delta_step=0,
                        max_depth=3, min_child_weight=1, missing=None, n_estimators=100,
                        n_jobs=1, nthread=None, objective='binary:logistic', random_state=0,
                        reg_alpha=0, reg_lambda=1, scale_pos_weight=1, seed=None,
                        silent=True, subsample=1)

In [22]: prob_predictions = xgboost.predict_proba(df_training_features)
         prob_predictions = np.append(prob_predictions, xgboost.predict_proba(df_validation_features))

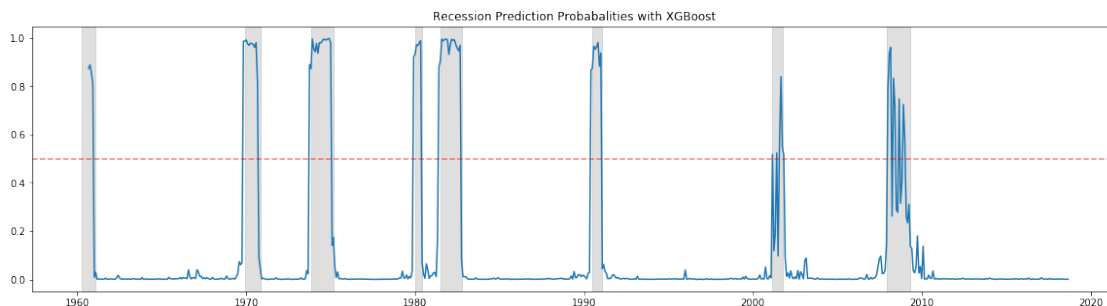
In [23]: sample_range = pd.date_range(start='9/1/1960', end='9/1/2018', freq='MS')

```

```

plt.figure(figsize=(20,5))
plt.plot(sample_range.to_series().values, prob_predictions[:,0])
for i in range(len(rec_spans)):
    plt.axvspan(rec_spans[i][0], rec_spans[i][len(rec_spans[i]) - 1], alpha=0.25, color='gray')
plt.axhline(y=0.5, color='r', ls='dashed', alpha = 0.5)
plt.title('Recession Prediction Probabilities with XGBoost')
mp.savefig('plot3.png', bbox_inches='tight')
plt.show()

```

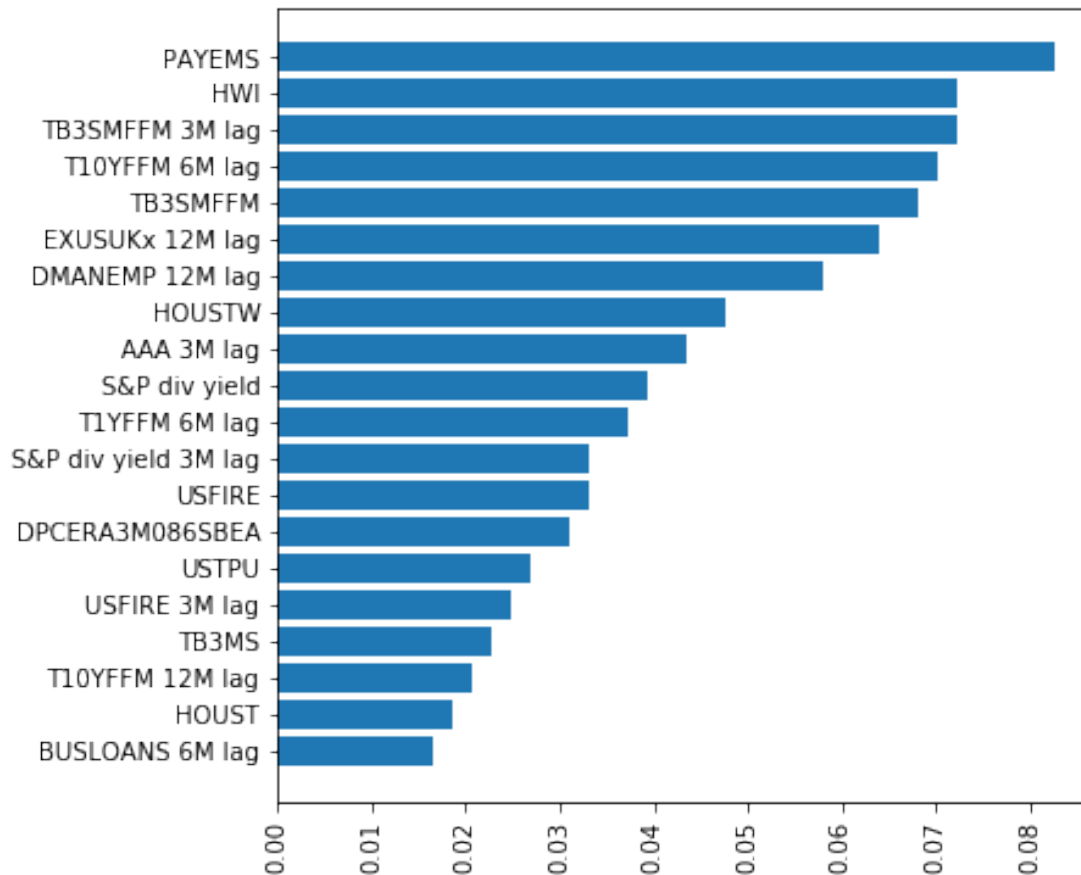


```

In [27]: # find feature importances
headers = df.drop(['Regime', 'Label', 'Date'], axis=1).columns.values.tolist()
xgboost_importances = pd.DataFrame(xgboost.feature_importances_, index = headers, columns = ['Relative Importance'])
_ = xgboost_importances.sort_values(by = ['Relative Importance'], ascending = False)
xgboost_importances = xgboost_importances[xgboost_importances['Relative Importance'] > 0.05]

# display importances in bar-chart and pie-chart
fig = plt.figure(figsize=(6,6))
plt.xticks(rotation='90')
plt.barh(y=np.arange(len(xgboost_importances)), width=xgboost_importances['Relative Importance'])
plt.gca().invert_yaxis()
mp.savefig('feature_importance.png', bbox_inches='tight')
plt.show()

```



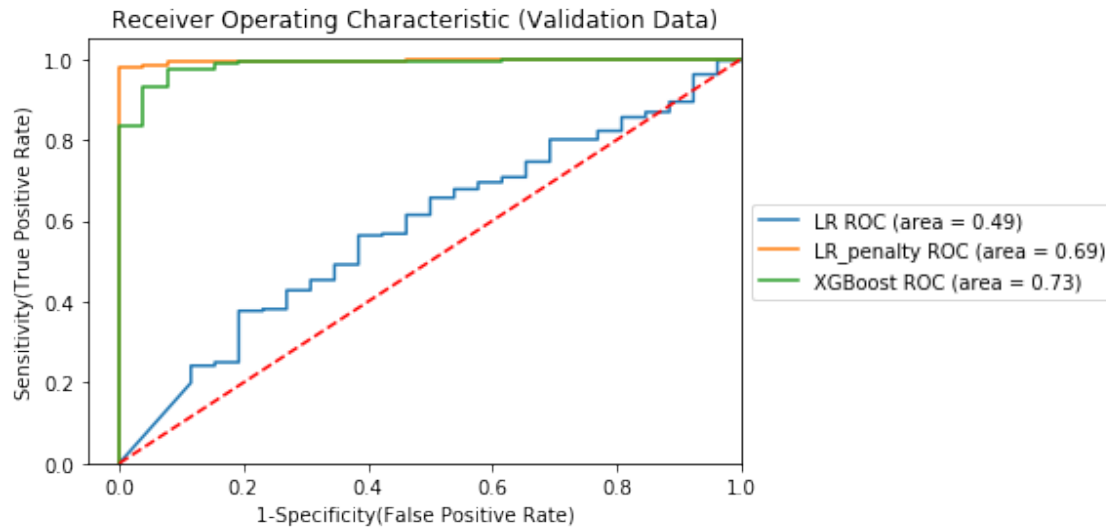
```
In [25]: fpr, tpr, thresholds = metrics.roc_curve(df_validation_targets, LR.predict_proba(df_validation_features))
auc = metrics.roc_auc_score(df_validation_targets, LR.predict(df_validation_features))
plt.plot(fpr, tpr, label='%s ROC (area = %0.2f)' % ('LR', auc))

fpr, tpr, thresholds = metrics.roc_curve(df_validation_targets, LR_penalty.predict_proba(df_validation_features))
auc = metrics.roc_auc_score(df_validation_targets, LR_penalty.predict(df_validation_features))
plt.plot(fpr, tpr, label='%s ROC (area = %0.2f)' % ('LR_penalty', auc))

fpr, tpr, thresholds = metrics.roc_curve(df_validation_targets, xgboost.predict_proba(df_validation_features))
auc = metrics.roc_auc_score(df_validation_targets, xgboost.predict(df_validation_features))
plt.plot(fpr, tpr, label='%s ROC (area = %0.2f)' % ('XGBoost', auc))

plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([-0.05, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('1-Specificity(False Positive Rate)')
plt.ylabel('Sensitivity(True Positive Rate)')
plt.title('Receiver Operating Characteristic (Validation Data)')
plt.legend(loc='center left', bbox_to_anchor=(1, 0.5))
```

```
mp.savefig('ROC1.png', bbox_inches='tight')
plt.show()
```



```
In [26]: fpr, tpr, thresholds = metrics.roc_curve(df_targets, LR.predict_proba(df_features)[:],
auc = metrics.roc_auc_score(df_targets,LR.predict(df_features))
plt.plot(fpr, tpr, label='%s ROC (area = %0.2f)' % ('LR', auc))

fpr, tpr, thresholds = metrics.roc_curve(df_targets, LR_penalty.predict_proba(df_features),
auc = metrics.roc_auc_score(df_targets,LR_penalty.predict(df_features))
plt.plot(fpr, tpr, label='%s ROC (area = %0.2f)' % ('LR_penalty', auc))

fpr, tpr, thresholds = metrics.roc_curve(df_targets, xgboost.predict_proba(df_features),
auc = metrics.roc_auc_score(df_targets,xgboost.predict(df_features))
plt.plot(fpr, tpr, label='%s ROC (area = %0.2f)' % ('XGBoost', auc))

plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([-0.05, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('1-Specificity(False Positive Rate)')
plt.ylabel('Sensitivity(True Positive Rate)')
plt.title('Receiver Operating Characteristic (Whole period)')
plt.legend(loc='center left', bbox_to_anchor=(1, 0.5))
mp.savefig('ROC2.png', bbox_inches='tight')
plt.show()
```

